

DFM Mimir v1: An Open HRM Delivering Frontier Performance at 1B Parameters Using Only Permissible Post-Training Data

Peter Schneider-Kamp^{1,2,*}, Jacob Nielsen^{1,2}, Gianluca Barmina¹, Kenneth Enevoldsen³, Lukas Galke Poech¹,

¹University of Southern Denmark ²Ordbogen A/S ³Aarhus University

Current large language model development relies on massive, often non-permissible datasets, creating a high barrier for researchers committed to open-source and ethically sourced data. We introduce Mimir v1, a 1-billion-parameter language model based on the Hierarchical Reasoning Model (HRM) architecture, that is trained from scratch and delivers highly competitive performance for English and sets a new state of the art for Danish using only permissible post-training data. Trained on a mixture of 161 datasets, Mimir v1 outperforms the original HRM-Text 1B and competes with larger frontier models like Qwen 3.5 4B and Gemma 4 E2B, tested across 20 benchmarks for English, Math & Code and Danish. The model is available on the Hugging Face Hub: <https://huggingface.co/danish-foundation-models/DFM-Mimir>



Date: 14 August 2026

Correspondence: Peter Schneider-Kamp via petersk@imada.sdu.dk




Figure 1: Aggregate results comparing DFM Mimir 1B against the HRM-Text 1B, Qwen 3.5 2B and Gemma 4 E2B, displaying highly competitive performance across 20 benchmarks.

1 Introduction

In recent years, Large Language Models (LLMs) have catalyzed a paradigm shift in artificial intelligence, characterized by rapid iterations and significant advancements in emergent capabilities. However, as noted by Wang et al. [2026], current development is largely driven by a “monolithic recipe” consisting of massive, multi-stage pipelines and training on exorbitant volumes of data. This approach not only necessitates vast computational resources for pre-training but also creates a prohibitive entry barrier for the broader community of researchers and practitioners. For national initiatives such as the Danish Foundation Models project¹ [Enevoldsen et al., 2023], which adheres to a philosophy of using exclusively permissible and, whenever possible, openly licensed data, training a capable LLM from scratch is often infeasible given the limited pool of high-quality data for a language such as Danish. Consequently, it has been challenging to provide fully permissible base models on LLM platforms to serve as foundations for post-training objectives.

To address these constraints, we employ the HRM-Text framework, which enables focusing on post-training data during the initial training phase, thereby facilitating the creation of a viable base model for the wider community. In this technical report, we present **Mimir v1**, a 1-billion-parameter hierarchical reasoning model (HRM) trained from scratch, utilizing the architecture proposed by Wang et al. [2026]. Mimir v1 is optimized for Danish and English tasks and has undergone instruction-tuning on a curated mixture of 161 datasets, comprising approximately 70.5 billion tokens per epoch. Furthermore, as certain datasets used in [Wang et al., 2026] do not align with DFM’s permissibility standards, we demonstrate the efficacy of synthetically generating “transplant datasets”, replacing non-permissible data with synthetically generated permissible variants.

Our results indicate that these synthetic alternatives achieve comparable or superior performance without compromising data rights, further underscoring the relevance of the HRM approach for low-resource linguistic domains, empowering communities of practitioners and researchers with small, capable and fully permissible models, with both low training and inference requirements.

2 Datasets

For training HRM-Mimir v1, we curated an extensive mix of data with different objectives, ranging from English and Danish instruction & knowledge to mathematics and agentic-style post-training data. Our mix draws from 161 datasets with almost all of them being freely available on the HuggingFace Hub. The corpus amounts to **70,479,308,606 tokens** per epoch. The full list of datasets and their sources is listed in Appendix A.

2.1 Category Distribution

We classify each dataset into one of eight functional categories based on its content and intended use. Table 1 provides the distribution. The three largest categories: Danish instruction & knowledge (22.07%), English instruction (19.26%), and selected datasets from the Sapien mixed collection (17.02%)², together account for over 58% of the corpus. Math & reasoning contributes a further 14.8%, bringing the combined share of the top four categories to 73%. The Danish instruction & knowledge category is the largest by token volume, driven primarily by lærebogen, a

¹<https://www.foundationmodels.dk>

²[sapieninc/HRM-Text-data-io-cleaned-20260515](https://huggingface.co/datasets/sapieninc/HRM-Text-data-io-cleaned-20260515)

Table 1: Token share and dataset count by functional category.

Category	Tokens/epoch	Share	Datasets	Avg. tokens/dataset
Danish instruction & knowledge	15.56B	22.07%	30	518.6M
English instruction	13.58B	19.26%	16	848.5M
Sapient mixed (Flan/Platypus)	12.00B	17.02%	71	169.0M
Math & reasoning	10.40B	14.76%	10	1,040.2M
Mimir synthetic	7.05B	10.00%	16	440.4M
Agentic & tool use	6.66B	9.46%	8	832.5M
Machine translation	3.50B	4.96%	5	699.0M
Science & summarization	1.74B	2.47%	5	348.5M
Total	70.48B	100%	161	

Danish instruction-following dataset contributing 8.32B tokens (11.8% of the corpus) at $4\times$ repetition, alongside dfm-dyna-instruct (3.54B) and synquid_wiki-instruct-da (0.99B). English instruction is anchored in Dolci [Team Olmo, 2025] (7.71B combined), Tulu 3 [Lambert et al., 2024] variants (1.57B), and Nemotron instruction-following (1.60B). The Sapient mixed category is dominated by a single large repository (11.92B, 16.9%), which bundles 107 sub-collections from Flan, Platypus, and tasksource. An additional 70 *Sapient-synth* transplant datasets contribute 75M tokens. These are synthetic recreations of English instruction tasks (Flan NIV2, Flan Dialog, Platypus, Tasksource) in generated-and-audited form, replacing original Sapient data that was non-compliant with the DFM philosophy. Math & reasoning is led by OpenMathInstruct-2 [Toshniwal et al., 2024] (6.60B, 9.4%), the third-largest single dataset overall, followed by AceReason-1.1-SFT [Liu et al., 2025] (1.95B) and verifiable reasoning traces (0.68B). The remaining four categories – synthetic, agentic & tool use, machine translation, and science & summarization – together contribute 19.0B tokens (27%).

2.2 Language Distribution

Table 2 shows the distribution of tokens by language. The corpus is predominantly English (68.5%), with Danish contributing 24.7% (6 out of 8 categories are entirely English) and bilingual Danish–English data a further 6.4%. A small fraction (0.2%) contains other bi-lingual translation data. The bilingual da+en category includes machine translation data and synthetic transformation datasets that pair Danish and English.

Table 2: Token share by language.

Language	Tokens/epoch	Share
English (en)	48.36B	68.62%
Danish (da)	17.44B	24.74%
Bilingual da+en	4.61B	6.54%
other	0.14B	0.20%

Table 3: Token share by data form.

	Reformatted	Curated + reformatted	Synthetic + audited	Tool-call formatted	Translated + audited	Agreement-supplied	Derived task
Tokens/epoch	46.49B	11.92B	7.81B	1.87B	1.59B	0.67B	0.13B
Share	65.96%	16.91%	11.08%	2.65%	2.26%	0.95%	0.18%

2.3 Data Processing

Datasets enter the corpus in seven distinct forms, reflecting the pipeline’s processing stages. We report the forms in Table 3. **Reformatted** datasets are existing Hugging Face repositories simply converted into the training format – this represents the default pathway for public data. **Curated + reformatted** applies to the Sapient mega-repository, whose 107 sub-collections were selected as a curated subset before reformatting. **Synthetic + audited** data is LLM-generated using Gemma4 31B and quality-audited before inclusion, with acceptance rates ranging from single digit percentages to high nineties for different categories. This includes transformations from high-quality English and Danish text corpora into span-filling, denoising, reordering, and continuation tasks. Here, we employ Common Pile [Kandpal et al., 2025] for English text corpora and Danish Dynaword [Enevoldsen et al., 2025] for Danish ones. We similarly generated and audited the dfm8-synthetic-* instruction datasets and Sapient-synth transplants. **Tool-call formatted** data incorporates native tool-calling structure for agentic training. **Translated + audited** covers the OpenHermes-based data and DA/EN translations that were repaired and audited. A small corpus of **agreement-supplied** data comes from Danish Foundation Model agreements (DBC, Lex.dk), where licensing does not permit public sharing. **Derived task** data is derived from an existing dataset to create a new task formulation.

The original Sapient training data³ consists mainly of Flan, Platypus, and tasksource sub-collections, which are themselves dominated by multiple-choice classification tasks — pick the correct option from A/B/C/D. This is a natural consequence of their source benchmarks: Flan NIV2, Platypus (Reclor, SciBench), and tasksource (PragmEval, Reclor) are structured around categorical selection. In this work, we shift the balance away from multiple-choice toward free-form generation, presenting a substantially harder task for the model to achieve exact match accuracy. The majority of the corpus (83% of tokens) comes from outside the Sapient collection, and the dominant non-Sapient categories are generally and inherently more generative:

- **English instruction** (13.58B, 19.3%): Dolci, Tulu 3, and Nemotron SFT mixtures are predominantly free-form instruction following.
- **Math & reasoning** (10.40B, 14.8%): OpenMathInstruct-2, AceReason, and reasoning traces all require the model to produce a free-form numerical or symbolic answer that is scored by exact match, not by selecting from options.
- **Danish instruction & knowledge** (15.56B, 22.1%): Laerebogen, dfm-dyna-instruct, and Danish QA/summarization datasets are open-ended generation tasks.

Together these three categories account for 39.54B tokens (56.1% of the corpus), nearly all scored on exact accuracy or free-form generation rather than multiple-choice selection. Even

³sapientinc/HRM-Text-data-io-cleaned-20260515

Table 4: Top 10 datasets by sampled tokens per epoch.

Dataset	Tokens/epoch	Share
sapientinc/HRM-Text-data-io-cleaned-20260515	11.92B	16.91%
danish-foundation-models/laerebogen	8.32B	11.81%
nvidia/OpenMathInstruct-2	6.60B	9.37%
nvidia/Nemotron-SFT-Agent-ic-v2	4.27B	6.06%
danish-foundation-models/dfm-dyna-instruct	3.54B	5.03%
allenai/Dolci-Instruct-SFT-No-Tools	3.49B	4.95%
schneiderkamplab/opus-da-en-permissive	2.90B	4.12%
allenai/Dolci-Instruct-SFT	2.24B	3.17%
nvidia/AceReason-1.1-SFT	1.95B	2.76%
allenai/big-reasoning-traces	1.66B	2.35%

within the Sapient-synth transplant datasets (70 datasets, 75M tokens), many original multiple-choice classification tasks were regenerated as open-ended generation or answer-generation formulation for example, ‘task590-amazonfood-summary-correction-classification’ and ‘task870-msmarco-answer-generation’ ask the model to produce a free-form answer rather than select from candidates.

This change of the training data composition means Mimir v1 is trained to generate answers rather than discriminate among options, aligning with evaluation suites that prioritise exact-match scoring (GSM8k, MATH, DROP) over multiple-choice accuracy (ARC-C, MMLU, Hellaswag).

2.4 Data Concentration and Repetition

Table 4 shows the top 10 datasets by sampled tokens per epoch. The corpus is highly concentrated: these ten datasets account for 66.5% of all tokens, and the top three alone for 38.1%, while the remaining 151 datasets contribute 33.5%. Two single sources exceed 10% each: the Sapient mega-repository (16.9%) and lærebogen (11.8%). The Dolci family contributes 5.73B tokens across two instruction datasets, the largest English instruction contribution after the cleaned Sapient corpus, and NVIDIA sources account for 12.82B tokens across the reasoning, math, and agentic datasets in the top 10. This concentration is partly by construction: Several datasets are sampled more than once per epoch. Lærebogen is repeated 4×, inflating 2.08B base tokens to 8.32B and making it the second-largest entry in the corpus. Eight small Danish datasets are repeated 10× to ensure sufficient coverage despite limited source data, and Dolci-Instruct-SFT-No-Tools is doubled to increase its English instruction contribution. The most heavily repeated dataset, kaenguruen (20×), is negligible in volume at 638K tokens.

3 Architecture

Mimir uses the Hierarchical Reasoning Model Text (HRM-Text) architecture with a hidden size of 1,536. The model has 12 attention heads per layer and a feed-forward expansion factor of 4. Hierarchical reasoning is configured with 2 H-cycles and 3 L-cycles, with truncated backpropagation limited to 5 steps and a warmup ratio of 0.2. Positional encoding uses Rotary Position Embedding (RoPE) with $\theta = 10,000$, and the model applies pre-norm layer normalisation with $\epsilon = 10^{-6}$. Table 5

Table 5: Model hyperparameters.

Hidden size	Layers	Half layers	Attn. heads	Expansion	H cycles	L cycles	BP max steps	BP warmup ratio	Pos. emb.	RoPE θ	Norm	Norm ϵ
1,536	32	true	12	4	2	3	5	0.2	RoPE	10,000	pre-norm	10^{-6}

summarises the model configuration.

4 Training

The Mimir model is trained from scratch using the Gemma-4 tokenizer [Gemma Team, 2026], whereas HRM-Text employs a custom one. Through the application of a chat template, the model learns the structural conventions and behavioural patterns characteristic of modern conversational AI. Mimir was trained with Fully Sharded Data Parallelism (FSDP) using bfloat16 as dtype for computation with fp32 used as gathering precision (conventional setup). We employ the AdamW optimizer [Loshchilov and Hutter, 2017] with a peak learning rate of 3×10^{-4} , 2,000-step linear warm-up, and a constant schedule thereafter (min ratio 1.0). We use a global batch size of 262,144 tokens with a gradient accumulation of 2 on 8 accelerators for a per-accelerator batch size of 16384. fitting 4 contexts of length 4096 each. Table 6 provides an overview of our hyperparameter values. Our openly available framework⁴ builds upon Sapient’s code for HRM-Text [Wang et al., 2026].

Table 6: Training hyperparameters.

LR	LR warmup	LR min ratio	Adam β_1/β_2	Weight decay	EMA decay	Global batch	Grad. accum. steps	Seed
3×10^{-4}	2,000	1.0	0.9 / 0.95	0.1	0.9999	262,144	2	0

We trained the model for 1.65M steps on 8 NVIDIA B200 GPUs with 180 GB HBM3 in just under 3 weeks with an average step time of just under 1.1 seconds.

5 Results

Tables 7, 8, and 9 show evaluation results across a broad range of English, Math & Code, and Danish benchmarks. We compare our Mimir model with other 1B-parameter models: HRM-Text, Qwen 3.5, Gemma 3 and OLMo2. Furthermore, we compare Mimir against models in the ranges 2-3B and 4-5B including Gemma 4 E2B with 5B total parameters (effective 2.3B). In Figure 1, we report the average score for English, Math & Code, and Danish benchmark suites.

On the English benchmarks, Mimir outperforms all considered competitors on BoolQ, Winogrande, and DROP. On Math & Code, Mimir leads across its weight-class for GSM8K and HumanEval, with Mimir being second overall on GSM8K and better than Qwen3.5 2B on HumanEval. On the Danish benchmarks, Mimir outperforms all competitors on grammatical tasks (DaLA, GEC), question-answering tasks (WikiQA), and is close to the best on Nordjylland News (N.News; summarization). On average, Mimir displays superior performance on the Danish benchmarks, is only 0.3 points behind Qwen 3.5 4B on English tasks, and only 3.8% behind SmolLM3 3B on Math & Code, which was the best tested conventional LLM of the 2–3B weight class. On Math & Code, Mimir yields a 36.7% improvement compared to HRM-Text (64.1 Mimir vs. 46.9 HRM-Text).

⁴<https://github.com/schneiderkamplab/HRM-Text>

Table 7: English benchmark results. Best scores in **bold**.

Model	BoolQ (Acc)	Winogrande (Acc)	Hellaswag (Acc)	MMLU (Acc)	ARC-C (Acc)	DROP (F1)	GovRep. (R1)	Avg.
~1B models								
Mimir 1B	87.8	73.5	67.3	57.5	81.6	83.1	32.0	69.0
HRM-Text 1B	87.5	70.4	60.4	58.7	82.2	78.1	25.4	66.1
Qwen 3.5 0.8B	69.8	48.9	37.0	51.5	68.4	45.2	32.5	50.5
Gemma 3 1B	62.4	49.1	30.6	37.5	43.5	7.0	29.5	37.1
OLMo 2 1B	67.2	51.0	42.4	41.6	48.1	12.4	37.7	42.9
2–3B models								
Qwen 3.5 2B	80.8	53.4	64.6	62.8	82.7	31.3	31.5	58.2
SmolLM3 3B	84.3	60.3	65.1	60.2	79.5	54.0	38.1	63.1
4–5B models								
Qwen 3.5 4B	87.0	70.0	83.2	75.8	92.9	48.0	27.9	69.3
Gemma 4 E2B	64.1	56.7	55.6	59.3	69.8	57.3	33.6	56.6
Gemma 4 E2B (think)	83.4	63.0	55.8	72.0	86.8	70.8	34.7	66.6

Table 8: Math & Code benchmark results. Best scores in **bold**.

Model	GSM8K (Acc)	MATH (Acc)	HumanEval (Acc)	Avg.
~1B models				
Mimir 1B	89.9	45.8	56.7	64.1
HRM-Text 1B	84.8	56.0	0.0	46.9
Qwen 3.5 0.8B	49.1	36.1	30.5	38.6
Gemma 3 1B	49.7	37.2	42.7	43.2
OLMo 2 1B	59.4	18.8	15.9	31.4
2–3B models				
Qwen 3.5 2B	73.7	55.7	47.6	59.0
SmolLM3 3B	80.0	62.2	61.6	67.9
4–5B models				
Qwen 3.5 4B	60.5	56.5	78.0	65.0
Gemma 4 E2B	88.3	64.2	73.8	75.4
Gemma 4 E2B (think)	90.3	49.1	72.0	70.5

Table 9: Danish benchmark results. Best scores in **bold**.

Model	Angry Tweets (Acc)	DaLA (F1)	GEC (EM)	PIQA (Acc)	Daisy (EM)	WikiQA (EM)	WMT (chrF)	N.News (chrF)	IFEval (Acc)	Hellaswag- DA (Acc)	Avg.
~1B models											
Mimir 1B	67.4	96.1	85.6	53.7	9.6	66.8	53.9	35.87	63.9	35.3	56.8
HRM-Text 1B	42.4	26.7	0.5	13.0	0.0	34.9	25.4	26.76	18.5	28.8	21.7
Qwen 3.5 0.8B	53.8	51.0	0.7	56.5	0.7	41.6	37.8	35.30	39.6	25.0	34.2
Gemma 3 1B	54.4	41.0	3.3	72.2	1.4	42.6	45.1	35.56	47.2	24.8	36.8
OLMo 2 1B	33.6	48.7	0.2	75.0	0.0	8.4	30.0	33.77	32.5	26.7	28.9
2–3B models											
Qwen 3.5 2B	61.6	36.4	8.0	25.0	2.5	49.4	45.6	34.85	56.1	24.7	34.4
SmolLM3 3B	63.2	33.5	3.3	51.9	2.2	0.3	37.3	35.98	49.8	40.1	31.7
4–5B models											
Qwen 3.5 4B	69.1	50.1	42.6	70.4	4.7	57.1	52.1	37.03	73.7	34.7	49.2
Gemma 4 E2B	64.6	56.7	36.9	46.3	5.6	44.1	55.2	35.67	75.5	25.6	44.6
Gemma 4 E2B (think)	67.7	66.8	23.4	63.9	5.1	59.3	56.0	36.30	81.2	39.0	49.9
8–9B models											
Munin-Apertus 8B	60.6	46.1	42.1	81.5	12.5	49.9	55.8	30.30	53.0	24.5	45.6
Munin-Mistral 8B	61.3	48.8	26.4	76.9	8.4	48.4	51.8	32.92	67.8	33.6	45.6
Munin-Qwen 9B	69.1	60.6	11.4	38.9	5.4	55.7	56.1	35.89	71.8	34.3	43.9

Evaluation Setup

All benchmarks were evaluated at temperature 0 (greedy decoding) with shuffle seed 4242 on full datasets. All models used vLLM-served [Kwon et al., 2023] endpoints with FlashInfer, with the exception of Mimir, which requires FlashAttention to correctly capture the PrefixLM and Gemma 4 chat template. We ran both vLLM with FlashAttention4 [Zadouri et al., 2026] and Hugging Face Transformers, obtaining comparable results up to numerical stability. For the ease of reproduction, we report the results from Hugging Face Transformers [Wolf et al., 2019]. Some English benchmarks use few-shot prompting, with the number of shots following Wang et al. [2026]’s evaluation config. All Danish tasks are 0-shot. MCQ tasks use `max_tokens=1`. Details in Appendix B. Gemma 4 was evaluated in two modes: non-thinking and thinking with vLLM flag `--reasoning-parser gemma4` to strip thinking tokens before scoring. Thinking requires ~500–650 tokens. All non-MCQ tasks (or whenever reasoning is enabled) use `max_tokens=2048`. All baseline evaluations were conducted via the Inspect AI Framework of the AI Security Institute [2024].

6 Conclusion, Limitations and Future Directions

In this technical report, we presented Mimir v1, a 1-billion-parameter language model that leverages the HRM-Text architecture to provide frontier-level performance using only permissible data, excluding data containing personal information or copyright infringement and including data that is either openly licensed, made available by agreement, or allowed by the European Union’s text and data mining exception for research institutions. By curating a diverse, large corpus, consisting of 70.5B tokens per epoch, including synthetic ‘transplant’ datasets, we have demonstrated that high-quality instruction following and reasoning capabilities can be achieved without relying on large-scale pre-training corpora or prohibited data sources. Mimir shows strong improvements over the baseline HRM-Text and remains competitive, or superior, to much larger models in several English and Danish benchmarks. Despite the generally highly competitive performance, Mimir v1

still lags behind Gemma 4 (5B, effective 2.3B) on the Math & Code domains, making room for improvement in future iterations.

Future work will focus on investigating scaling behavior of HRM models like Mimir. Moreover, even though Mimir is trained with a Gemma 4 chat template from scratch, the capabilities as an assistant are still limited compared to the state of the art. This calls for future work in this capacity, including reinforcement learning, which is yet unexplored for this architecture. Lastly, we will continue developing the dataset to achieve full openness regarding licensing and further improved model performance.

7 Contributors

We list contributions according to the Contributor Roles Taxonomy (CRediT)⁵:

Peter Schneider-Kamp *Conceptualization, Data curation, Formal analysis, Funding acquisition, Investigation, Methodology, Project administration, Resources, Software, Supervision, Validation, Visualization, Writing – original draft, Writing – review & editing*

Jacob Nielsen *Formal analysis, Investigation, Methodology, Software, Validation, Visualization, Writing – original draft, Writing – review & editing*

Lukas Galke Poech *Formal analysis, Investigation, Methodology, Software, Supervision, Validation, Visualization, Writing – review & editing*

Gianluca Barmina *Data curation, Investigation, Writing – review & editing*

Kenneth Enevoldsen *Data curation, Resources, Writing – review & editing*

Additional Contributors

Mogens Henrik From *Investigation, Project administration, Software*

Andrea Blasi Núñez *Formal analysis, Software, Visualization*

Annemette Brok Pirchert *Formal analysis, Software, Visualization*

Stine Lyngsø Beltoft *Data curation, Resources*

Torben Blach *Project administration, Funding acquisition*

Sofie Helene Bruun *Data curation, Resources*

Oliver Kinch *Data curation, Resources*

Rasmus Larsen *Data curation, Resources*

Dan Saatrup Smart *Data curation, Resources*

Kristoffer Laigaard Nielbo *Funding acquisition*

⁵<https://credit.niso.org>

8 Acknowledgements

This work originated within OdenseNLP⁶ and was supported by the Danish Foundation Models project – a collaboration between the University of Southern Denmark, Aarhus University, the Alexandra Institute, and the University of Copenhagen – funded by the Ministry of Science, Higher Education and Digital Affairs. We thank all members of the Danish Foundation Models project and the OdenseNLP research group for many fruitful discussions and their continued support.

References

- UK AI Security Institute. Inspect AI: Framework for Large Language Model Evaluations, 2024. URL https://github.com/UKGovernmentBEIS/inspect_ai.
- Kenneth Enevoldsen, Lasse Hansen, Dan S. Nielsen, Rasmus A. F. Egebæk, Søren V. Holm, Martin C. Nielsen, Martin Bernstorff, Rasmus Larsen, Peter B. Jørgensen, Malte Højmark-Bertelsen, Peter B. Vahlstrup, Per Møldrup-Dalum, and Kristoffer Nielbo. Danish foundation models, 2023. URL <https://arxiv.org/abs/2311.07264>.
- Kenneth Enevoldsen, Kristian Nørgaard Jensen, Jan Kostkan, Balázs Szabó, Márton Kardos, Kirsten Vad, Johan Heinsen, Andrea Blasi Núñez, Gianluca Barmina, Jacob Nielsen, Rasmus Larsen, Peter Vahlstrup, Per Møldrup Dalum, Desmond Elliott, Lukas Galke, Peter Schneider-Kamp, and Kristoffer Nielbo. Dynaword: From one-shot to continuously developed datasets, 2025.
- Gemma Team. Gemma 4 technical report, 2026. URL <https://arxiv.org/abs/2607.02770>.
- Nikhil Kandpal, Brian Lester, Colin Raffel, Sebastian Majstorovic, Stella Biderman, Baber Abbasi, Luca Soldaini, Enrico Shippole, A. Feder Cooper, Aviya Skowron, John Kirchenbauer, Shayne Longpre, Lintang Sutawika, Alon Albalak, Zhenlin Xu, Guilherme Penedo, Loubna Ben Allal, Elie Bakouch, John David Pressman, Honglu Fan, Dashiell Stander, Guangyu Song, Aaron Gokaslan, Tom Goldstein, Brian R. Bartoldson, Bhavya Kailkhura, and Tyler Murray. The common pile v0.1: An 8tb dataset of public domain and openly licensed text, 2025. URL <https://arxiv.org/abs/2506.05209>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*, 2023.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, et al. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*, 2024.
- Zihan Liu, Zhuolin Yang, Yang Chen, Chankyu Lee, Mohammad Shoeybi, Bryan Catanzaro, and Wei Ping. Acereason-nemotron 1.1: Advancing math and code reasoning through sft and rl synergy. *arXiv preprint arXiv:2506.13284*, 2025.

⁶<https://odensenlp.github.io>

Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.

Team Olmo. Olmo 3, 2025. URL <https://arxiv.org/abs/2512.13961>.

Shubham Toshniwal, Wei Du, Ivan Moshkov, Branislav Kisacanin, Alexan Ayrapetyan, and Igor Gitman. Openmathinstruct-2: Accelerating ai for math with massive open-source instruction data. *arXiv preprint arXiv:2410.01560*, 2024.

Guan Wang, Changling Liu, Chenyu Wang, Cai Zhou, Yuhao Sun, Yifei Wu, Shuai Zhen, Luca Scimeca, and Yasin Abbasi Yadkori. Hrm-text: Efficient pretraining beyond scaling. *arXiv preprint arXiv:2605.20613*, 2026.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. Huggingface’s transformers: State-of-the-art natural language processing. *arXiv preprint arXiv:1910.03771*, 2019.

Ted Zadouri, Markus Hoehnerbach, Jay Shah, Vijay Thakkar, and Tri Dao. Flashattention-4: Algorithm and kernel pipelining co-design for asymmetric hardware scaling. *Proceedings of Machine Learning and Systems*, 8:912–926, 2026.

A List of Training Datasets

Table 10 lists all the datasets with their corresponding HuggingFace identifier, the form type, how many tokens they represent, and the share in percentages.

Table 10: All 161 datasets in the Mimir corpus, sorted by sampled tokens per epoch.

#	Dataset	Form	Tokens/epoch	Share
1.	sapientinc/HRM-Text-data-io-cleaned-20260515	Curated + reformatted	11.92B	16.91%
2.	danish-foundation-models/laerebogen	Reformatted	8.32B	11.81%
3.	nvidia/OpenMathInstruct-2	Reformatted	6.60B	9.37%
4.	nvidia/Nemotron-SFT-Agent-ic-v2	Reformatted	4.27B	6.06%
5.	danish-foundation-models/dfm-dyna-instruct	Reformatted	3.54B	5.03%
6.	allenai/Dolci-Instruct-SFT-No-Tools	Reformatted	3.49B	4.95%
7.	schneiderkamplab/opus-da-en-permissive	Reformatted	2.90B	4.12%
8.	allenai/Dolci-Instruct-SFT	Reformatted	2.24B	3.17%
9.	nvidia/AceReason-1.1-SFT	Reformatted	1.95B	2.76%
10.	allenai/big-reasoning-traces	Reformatted	1.66B	2.35%
11.	allenai/Dolci-Instruct-SFT-Tool-Use	Tool-call formatted	1.61B	2.29%
12.	nvidia/Nemotron-SFT-Instruction-Following-Chat-v2	Reformatted	1.60B	2.27%
13.	allenai/tulu-3-sft-mixture	Reformatted	1.57B	2.23%
14.	laion/Scientific-Summaries	Reformatted	1.27B	1.80%
15.	schneiderkamplab/common-pile-prefix-continuation	Synthetic + audited	1.15B	1.64%
16.	synquid/wiki-instruct-da	Reformatted	988M	1.40%
17.	schneiderkamplab/dfm8-openhermes-da	Translated + audited	922M	1.31%
18.	schneiderkamplab/common-pile-denoising	Synthetic + audited	884M	1.25%
19.	allenai/tulu-v2-sft-mixture	Reformatted	840M	1.19%
20.	schneiderkamplab/common-pile-span-filling	Synthetic + audited	822M	1.17%
21.	schneiderkamplab/dfm8-openhermes-en	Translated + audited	672M	0.95366%
22.	allenai/tulu-v2-sft-long-mixture	Reformatted	605M	0.85902%
23.	allenai/verifiable-reasoning-filtered-gpt-41	Reformatted	605M	0.85798%
24.	open-thoughts/OpenThoughts2-1M	Reformatted	534M	0.75834%

continued from previous page

#	Dataset	Form	Tokens/epoch	Share
25.	schneiderkamplab/dfm8-synthetic-native-tool-calling	Synthetic + audited	513M	0.728%
26.	schneiderkamplab/transformations-danish-danish	Synthetic + audited	496M	0.70403%
27.	schneiderkamplab/transformations-english-english	Synthetic + audited	471M	0.66842%
28.	schneiderkamplab/transformations-danish-english	Synthetic + audited	469M	0.66529%
29.	schneiderkamplab/transformations-english-danish	Synthetic + audited	418M	0.59274%
30.	nvidia/Nemotron-SFT-Multilingual-v1	Reformatted	413M	0.58565%
31.	schneiderkamplab/dfm8-synthetic-danish-summarization-rewrite-controls	Synthetic + audited	396M	0.56181%
32.	MegaScience/TextbookReasoning	Reformatted	374M	0.5311%
33.	schneiderkamplab/dfm8-synthetic-multiturn-danish-english-chat	Synthetic + audited	366M	0.51997%
34.	DBC (agreement-supplied)	Agreement-supplied	356M	0.50529%
35.	schneiderkamplab/danish-dynaword-denoising	Synthetic + audited	323M	0.45861%
36.	Lex.dk articles	Agreement-supplied	313M	0.44467%
37.	oliverkinch/machine-translation-da-en	Reformatted	284M	0.40348%
38.	schneiderkamplab/danish-dynaword-prefix-continuation	Synthetic + audited	252M	0.35728%
39.	schneiderkamplab/danish-dynaword-span-filling	Synthetic + audited	251M	0.35593%
40.	schneiderkamplab/dfm8-synthetic-code-debugging	Synthetic + audited	247M	0.34996%
41.	GEM/wiki_cat_sum	Reformatted	205M	0.29117%
42.	allenai/tulu-3-sft-personas-math	Reformatted	204M	0.28922%
43.	giannor/gec_dala_tv2r_it	Reformatted	193M	0.27377%
44.	synquid/wildchat-100k-qwen-messages	Reformatted	190M	0.2694%
45.	schneiderkamplab/dfm8-synthetic-strict-math-answer-contract	Synthetic + audited	173M	0.24605%
46.	schneiderkamplab/dfm8-synthetic-constrained-format-following	Synthetic + audited	171M	0.24275%
47.	oliverkinch/danish-summarization	Reformatted	168M	0.23888%
48.	schneiderkamplab/common-pile-paragraph-reordering	Synthetic + audited	167M	0.23653%
49.	schneiderkamplab/danish-dynaword-paragraph-reordering	Synthetic + audited	160M	0.22769%

continued from previous page

#	Dataset	Form	Tokens/epoch	Share
50.	glaiveai/glaive-function-calling-v2	Tool-call format- ted	156M	0.22066%
51.	oliverkinch/machine-translation-da-ar	Reformatted	140M	0.19802%
52.	allenai/SciRIFF-train-mix	Reformatted	132M	0.18771%
53.	common-pile/arxiv_papers_filtered	Derived task	130M	0.18386%
54.	oliverkinch/da-instruct-dynaword	Reformatted	103M	0.14592%
55.	synquid/translation-100k	Reformatted	98.2M	0.13938%
56.	kobprof/skolegpt-instruct	Reformatted	86.4M	0.12259%
57.	allenai/verifiable-reasoning-filtered-o4-mini	Reformatted	80.2M	0.11378%
58.	oliverkinch/tidsskrift-dk-bt	Reformatted	76.7M	0.10878%
59.	allenai/open_math_2.50k_r1-original	Reformatted	71.8M	0.1019%
60.	oliverkinch/machine-translation-da-uk	Reformatted	69.5M	0.09859%
61.	giannor/dala_tv2r_it	Reformatted	68.5M	0.09723%
62.	Salesforce/xlam-function-calling-60k	Tool-call format- ted	67.0M	0.09502%
63.	oliverkinch/danish-qa	Reformatted	57.5M	0.08153%
64.	oliverkinch/dst-table-prompts-bt	Reformatted	51.1M	0.07248%
65.	danish-foundation- models/ai_arena_udtraek	Reformatted	45.7M	0.06483%
66.	oliverkinch/multi-wiki-qa-high-quality-subset	Reformatted	41.9M	0.05947%
67.	oliverkinch/dynaword-bt	Reformatted	34.3M	0.04867%
68.	HuggingFaceH4/no_robots	Reformatted	31.4M	0.04459%
69.	oliverkinch/da-instruct-dynaword-contemporary-hq	Reformatted	25.6M	0.03626%
70.	Team-ACE/ToolACE	Tool-call format- ted	25.2M	0.03579%
71.	oliverkinch/da-instruct-dynaword-hq	Reformatted	24.1M	0.03414%
72.	oliverkinch/danish-university-portals-bt	Reformatted	21.8M	0.03088%
73.	allenai/tulu-3-sft-personas-algebra	Reformatted	21.0M	0.02982%
74.	allenai/IF_sft_data_verified	Reformatted	19.9M	0.02829%
75.	oliverkinch/autodata-da-sft	Reformatted	18.9M	0.02677%
76.	oliverkinch/danmarks-statistik-bt	Reformatted	18.9M	0.02676%
77.	synquid/danish-verifiable-reasoning	Reformatted	18.1M	0.02573%
78.	oliverkinch/da-instruct-dynaword-contemporary	Reformatted	17.6M	0.02491%
79.	oliverkinch/eur-lex-bt	Reformatted	16.7M	0.02369%
80.	allenai/tulu-3-sft-personas-code	Reformatted	14.6M	0.02077%
81.	oliverkinch/instruct-bt	Reformatted	13.5M	0.0191%
82.	schneiderkamplab/sapient-synth-flan-dialog-fsopt-data-qrecc	Synthetic + au- dited	13.1M	0.0186%

continued from previous page

#	Dataset	Form	Tokens/epoch	Share
83.	synquid/ifbench-train	Reformatted	12.7M	0.018%
84.	allenai/tulu-3-sft-personas-instruction-following	Reformatted	11.5M	0.01634%
85.	schneiderkamplab/sapient-synth-flan-flan-fsopt-data-aeslc-1.0.0	Synthetic + audited	10.2M	0.01441%
86.	synquid/mt-da-deepseek	Reformatted	8.9M	0.01261%
87.	allenai/RLVR-MATH	Reformatted	8.1M	0.01145%
88.	allenai/RLVR-GSM	Reformatted	6.8M	0.00968%
89.	allenai/Dolci-Instruct-SFT-Tool-Use-SA	Tool-call formatted	5.8M	0.00825%
90.	schneiderkamplab/sapient-synth-flan-flan-fsnoopt-data-aeslc-1.0.0	Synthetic + audited	4.8M	0.0068%
91.	ccdv/govreport-summarization	Reformatted	4.4M	0.00626%
92.	schneiderkamplab/sapient-synth-flan-flan-fsopt-data-opinion-abstracts-rotten-tomatoes	Synthetic + audited	4.3M	0.00612%
93.	schneiderkamplab/sapient-synth-flan-dialog-fsopt-data-qrecc-ii	Synthetic + audited	2.9M	0.00413%
94.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task589-amazonfood-summary-text-generation	Synthetic + audited	2.8M	0.00399%
95.	oliverkinch/eur-lex-sum-instruct	Reformatted	2.6M	0.00373%
96.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task590-amazonfood-summary-correction-classification	Synthetic + audited	2.5M	0.0036%
97.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task618-amazonreview-summary-text-generation	Synthetic + audited	2.2M	0.00309%
98.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1309-amazonreview-summary-classification	Synthetic + audited	1.9M	0.00276%
99.	schneiderkamplab/sapient-synth-flan-flan-zsnoopt-data-aeslc-1.0.0	Synthetic + audited	1.9M	0.00271%
100.	schneiderkamplab/sapient-synth-flan-flan-fsnoopt-data-opinion-abstracts-rotten-tomatoes	Synthetic + audited	1.7M	0.00247%
101.	schneiderkamplab/sapient-synth-flan-flan-zsopt-data-aeslc-1.0.0	Synthetic + audited	1.7M	0.00245%
102.	schneiderkamplab/sapient-synth-flan-dialog-zsopt-data-qrecc	Synthetic + audited	1.7M	0.00245%

continued from previous page

#	Dataset	Form	Tokens/epoch	Share
103.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1375-newscomm-translation	Synthetic + audited	1.2M	0.00166%
104.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1373-newscomm-translation	Synthetic + audited	1.1M	0.0015%
105.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1376-newscomm-translation	Synthetic + audited	974K	0.00138%
106.	schneiderkamplab/sapient-synth-platypus-reclor	Synthetic + audited	922K	0.00131%
107.	schneiderkamplab/sapient-synth-tasksources-reclor	Synthetic + audited	901K	0.00128%
108.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task902-deceptive-opinion-spam-classification	Synthetic + audited	878K	0.00125%
109.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task634-allegro-reviews-classification	Synthetic + audited	816K	0.00116%
110.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task635-allegro-reviews-answer-generation	Synthetic + audited	816K	0.00116%
111.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task590-amazonfood-summary-correction-classification	Synthetic + audited	811K	0.00115%
112.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1309-amazonreview-summary-classification	Synthetic + audited	795K	0.00113%
113.	schneiderkamplab/sapient-synth-flan-flan-fsopt-data-opinion-abstracts-idebate	Synthetic + audited	746K	0.00106%
114.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1370-newscomm-classification	Synthetic + audited	721K	0.00102%
115.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task589-amazonfood-summary-text-generation	Synthetic + audited	705K	0.001%
116.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1371-newscomm-translation	Synthetic + audited	703K	0.001%
117.	oliverkinch/doab-da-bt	Reformatted	694K	0.00098%

continued from previous page

#	Dataset	Form	Tokens/epoch	Share
118.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1377-newscomm-translation	Synthetic + audited	679K	0.00096%
119.	danish-foundation-models/kaenguruen	Reformatted	638K	0.0009%
120.	schneiderkamplab/sapient-synth-flan-flan-zsopt-data-opinion-abstracts-rotten-tomatoes	Synthetic + audited	622K	0.00088%
121.	schneiderkamplab/sapient-synth-flan-flan-zsnoopt-data-opinion-abstracts-rotten-tomatoes	Synthetic + audited	616K	0.00087%
122.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task870-msmarco-answer-generation	Synthetic + audited	599K	0.00085%
123.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task1374-newscomm-translation	Synthetic + audited	599K	0.00085%
124.	schneiderkamplab/sapient-synth-tasksources-pragmeval-sarcasm	Synthetic + audited	542K	0.00077%
125.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task903-deceptive-opinion-spam-classification	Synthetic + audited	535K	0.00076%
126.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task618-amazonreview-summary-text-generation	Synthetic + audited	532K	0.00076%
127.	schneiderkamplab/sapient-synth-flan-flan-fsnoopt-data-opinion-abstracts-idebate	Synthetic + audited	527K	0.00075%
128.	schneiderkamplab/sapient-synth-flan-dialog-zsopt-data-qrecc-ii	Synthetic + audited	513K	0.00073%
129.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task265-paper-reviews-language-identification	Synthetic + audited	395K	0.00056%
130.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task635-allegro-reviews-answer-generation	Synthetic + audited	320K	0.00045%
131.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task909-dialogre-prevalent-speakers	Synthetic + audited	296K	0.00042%
132.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1376-newscomm-translation	Synthetic + audited	274K	0.00039%

continued from previous page

#	Dataset	Form		Tokens/epoch	Share
133.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1373-newscomm-translation	Synthetic	+ au-dited	268K	0.00038%
134.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1375-newscomm-translation	Synthetic	+ au-dited	265K	0.00038%
135.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task266-paper-reviews-reviewer-perspective	Synthetic	+ au-dited	236K	0.00033%
136.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task906-dialogre-identify-names	Synthetic	+ au-dited	235K	0.00033%
137.	schneiderkamplab/sapient-synth-flan-flan-zsnoopt-data-opinion-abstracts-idebate	Synthetic	+ au-dited	223K	0.00032%
138.	schneiderkamplab/sapient-synth-flan-flan-zsopt-data-opinion-abstracts-idebate	Synthetic	+ au-dited	223K	0.00032%
139.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task634-allegro-reviews-classification	Synthetic	+ au-dited	215K	0.0003%
140.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task870-msmarco-answer-generation	Synthetic	+ au-dited	212K	0.0003%
141.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task902-deceptive-opinion-spam-classification	Synthetic	+ au-dited	198K	0.00028%
142.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task672-amazon-yelp-summarization	Synthetic	+ au-dited	198K	0.00028%
143.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1377-newscomm-translation	Synthetic	+ au-dited	182K	0.00026%
144.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1374-newscomm-translation	Synthetic	+ au-dited	174K	0.00025%
145.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1371-newscomm-translation	Synthetic	+ au-dited	171K	0.00024%
146.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task265-paper-reviews-language-identification	Synthetic	+ au-dited	154K	0.00022%

continued from previous page

#	Dataset	Form		Tokens/epoch	Share
147.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task903-deceptive-opinion-spam-classification	Synthetic	+ au-dited	149K	0.00021%
148.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task264-paper-reviews-accept-reject	Synthetic	+ au-dited	140K	0.0002%
149.	schneiderkamplab/sapient-synth-platypus-scbench	Synthetic	+ au-dited	104K	0.00015%
150.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task672-amazon-yelp-summarization	Synthetic	+ au-dited	85.3K	0.00012%
151.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task909-dialogre-prevalent-speakers	Synthetic	+ au-dited	82.0K	0.00012%
152.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task266-paper-reviews-reviewer-perspective	Synthetic	+ au-dited	71.9K	0.0001%
153.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task871-msmarco-question-generation	Synthetic	+ au-dited	71.1K	0.0001%
154.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task908-dialogre-identify-familial-relationships	Synthetic	+ au-dited	68.4K	0.0001%
155.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task906-dialogre-identify-names	Synthetic	+ au-dited	57.5K	0.00008%
156.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task264-paper-reviews-accept-reject	Synthetic	+ au-dited	49.3K	0.00007%
157.	schneiderkamplab/sapient-synth-flan-niv2-fsopt-data-task907-dialogre-identify-relationships	Synthetic	+ au-dited	46.0K	0.00007%
158.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task871-msmarco-question-generation	Synthetic	+ au-dited	35.6K	0.00005%
159.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task908-dialogre-identify-familial-relationships	Synthetic	+ au-dited	17.3K	0.00002%
160.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task907-dialogre-identify-relationships	Synthetic	+ au-dited	13.4K	0.00002%

continued from previous page

#	Dataset	Form	Tokens/epoch	Share
161.	schneiderkamplab/sapient-synth-flan-niv2-zsopt-data-task1370-newscomm-classification	Synthetic + audited	1.9K	<0.00001%

B List of Evaluation Datasets

Table 11 lists all considered evaluation datasets with their precise source as HuggingFace identifier, and the number of in-context shots provided to the models during evaluation.

Table 11: Overview of benchmarks and their configuration. All benchmarks are evaluated with *temperature* 0 and a fixed seed.

Task	Dataset (split)	N-shots
English		
BoolQ	google/boolq	5
Winogrande	allenai/winogrande	5
Hellaswag	Rowan/hellaswag	10
MMLU	cais/mmlu	5
ARC-C	allenai/ai2_arc	25
DROP	EleutherAI/drop	3
GovReport	ccdv/govreport	0
Math&Code		
GSM8K	openai/gsm8k	0
MATH	EleutherAI/hendrycks_math	0
HumanEval	human_eval (164 problems)	0
Danish		
Angry Tweets	DDSC/angry-tweets	0
DaLA	giannor/dala	0
GEC-DaLA	giannor/dala_gen.v3	0
PIQA-da	local JSON	0
Daisy	schneiderkamplab/SDU-Daisy	0
Multi Wiki QA	oliverkinch/multi-wiki-qa	0
WMT24++ EN-DA	synquid/wmt24pp	0
Nordj. News Summ.	alexandrinst/nordj-news	0
IFEval-Da	danish-foundation-models/ifeval-da	0
Hellaswag-da	EuroEval	0